

# 1. Demonstrate data cleaning – missing values

```
library(tidyverse)
x <- sample(1:21, 20, replace = TRUE)
y <- sample(1:10, 20, replace = TRUE)
for(i in 1:20)
{
  a <- x[i]
  b <- y[i]
  mtcars[a, b] = NA
}
which(is.na(mtcars))
sum(is.na(mtcars))
na.exclude(mtcars)
view(mtcars)
dispna <- apply(mtcars["disp"], 2, mean, na.rm=TRUE)
view(dispna)
newcars <- mtcars %>%
  mutate(displacement = ifelse(is.na(displacement), dispna, displacement), )
view(newcars)
```

## Output

```
> which(is.na(mtcars))
[1] 1 10 33 37 42 48 66 69 73 76 77 85 101 105 112 115 116 136 149 16
2 170 171
[23] 174 175 193 194 196 203 206 213 239 245 261 290 298 305

> sum(is.na(mtcars))
[1] 36
```

> na.exclude(mtcars)

mpg cyl disp hp drat wt qsec vs am gear carb

Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

## 2. Implement data normalization (min-max, z-score)

```
arr <- c(9.5, 6.2, 8.9, 15.2, 20.0, 10.1, 5.4, 3.2, 1.0, 22.5, 10.0, 16.0)
```

```
#min-max
```

```
minarr <- min(arr)
```

```
maxarr <- max(arr)
```

```
arr2 <- arr
```

```
for (i in 1:12){
```

```
  arr2[i] = round((arr[i]-minarr)/(maxarr-minarr))
```

```
}
```

```
print(arr2)
```

```
#z-score
```

```
meanarr <- mean(arr)
```

```
sdarr <- sd(arr)
```

```
for (i in 1:12){
```

```
  arr2[i] = round((arr[i]-meanarr)/sdarr, 2)
```

```
}
```

```
print(arr2)
```

### **Output:**

```
> print(arr2)
```

```
[1] 0 0 0 1 1 0 0 0 0 1 0 1
```

```
>
```

```
> #z-score
```

```
> meanarr <- mean(arr)
```

```
> sdarr <- sd(arr)
```

```
> for (i in 1:12){
```

```
+ arr2[i] = round((arr[i]-meanarr)/sdarr, 2)
```

```
+ }
```

```
> print(arr2)
```

```
[1] -0.18 -0.68 -0.27 0.69 1.42 -0.09 -0.80 -1.13 -1.47 1.79 -0.10 0.81
```

### 3. Implement attribute subset selection for data reduction

```
# Install and load the leaps package (only needs to be installed once)
if (!require(leaps)) install.packages("leaps")
library(leaps)
View(as.data.frame(Titanic))
Titanic <- as.data.frame(Titanic)
sum(is.na(Titanic))
Titanic <- na.omit(Titanic)
dim(Titanic)
fwd <- regsubsets(Freq ~ ., data = Titanic, nvmax = 19, method = "forward")
bwd <- regsubsets(Freq ~ ., data = Titanic, nvmax = 19, method = "backward")
full <- regsubsets(Freq ~ ., data = Titanic, nvmax = 19)
summary(fwd)
summary(bwd)
summary(full)
coef(fwd, 3)
coef(bwd, 3)
coef(full, 3)
```

#### **Output:**

```
> summary(fwd)
```

Subset selection object

Call: regsubsets.formula(Freq ~ ., data = Titanic, nvmax = 19, method = "forward")

6 Variables (and intercept)

Forced in Forced out

Class2nd	FALSE	FALSE
Class3rd	FALSE	FALSE
ClassCrew	FALSE	FALSE
SexFemale	FALSE	FALSE
AgeAdult	FALSE	FALSE
SurvivedYes	FALSE	FALSE

1 subsets of each size up to 6

Selection Algorithm:

forward

	Class	2nd Class	3rd Class	Crew	Sex	Female	Age	Adult	Survived	Yes
1 ( 1 )	" "	" "	" "	" "	"*"	" "				
2 ( 1 )	" "	" "	" "	"*"	"*"	" "			>	s
3 ( 1 )	" "	" "	" "	"*"	"*"	"*"				u
4 ( 1 )	" "	" "	"*"	"*"	"*"	"*"				m
5 ( 1 )	" "	"*"	"*"	"*"	"*"	"*"				m
6 ( 1 )	"*"	"*"	"*"	"*"	"*"	"*"				a r

y(bwd) Subset

selection object

Call: regsubsets.formula(Freq ~ ., data = Titanic, nvmax = 19, method = "backward")

6 Variables (and intercept)

Forced in Forced out

Class2nd	FALSE	FALSE
Class3rd	FALSE	FALSE
ClassCrew	FALSE	FALSE
SexFemale	FALSE	FALSE
AgeAdult	FALSE	FALSE
SurvivedYes	FALSE	FALSE

1 subsets of each size up to 6

Selection Algorithm: backward

	Class	2nd Class	3rd Class	Crew	Sex	Female	Age	Adult	Survived	Yes
1 ( 1 )	" "	" "	" "	" "	"*"	" "				
2 ( 1 )	" "	" "	" "	"*"	"*"	" "				
3 ( 1 )	" "	" "	" "	"*"	"*"	"*"				
4 ( 1 )	" "	" "	"*"	"*"	"*"	"*"				
5 ( 1 )	" "	"*"	"*"	"*"	"*"	"*"				
6 ( 1 )	"*"	"*"	"*"	"*"	"*"	"*"				

> summary(full)

Subset selection object

Call: regsubsets.formula(Freq ~ ., data = Titanic, nvmax = 19)

6 Variables (and intercept)

Forced in Forced out

Class2nd	FALSE	FALSE
Class3rd	FALSE	FALSE
ClassCrew	FALSE	FALSE
SexFemale	FALSE	FALSE

AgeAdult FALSE FALSE  
SurvivedYes FALSE FALSE 1  
subsets of each size up to 6

Selection Algorithm: exhaustive

	Class2nd	Class3rd	ClassCrew	SexFemale	AgeAdult	SurvivedYes
1 ( 1 )	" "	" "	" "	" "	"*"	" "
2 ( 1 )	" "	" "	" "	"*"	"*"	" "
3 ( 1 )	" "	" "	" "	"*"	"*"	"*"
4 ( 1 )	" "	" "	"*"	"*"	"*"	"*"
5 ( 1 )	" "	"*"	"*"	"*"	"*"	"*"
6 ( 1 )	"*"	"*"	"*"	"*"	"*"	"*"

>

> coef(fwd, 3)

(Intercept)	SexFemale	AgeAdult	SurvivedYes
70.5625	-78.8125	123.9375	-48.6875

> coef(bwd, 3)

(Intercept)	SexFemale	AgeAdult	SurvivedYes
70.5625	-78.8125	123.9375	-48.6875

> coef(full, 3)

(Intercept)	SexFemale	AgeAdult	SurvivedYes
70.5625	-78.8125	123.9375	-48.6875

>

## 4. Demonstrate outlier detection

```
install.packages("tidyr") # For drop_na function
library(tidyr)
set.seed(123)
day <- data.frame(
  temp = rnorm(100, mean = 20, sd = 5),    # Random temperatures
  hum = rnorm(100, mean = 60, sd = 10),    # Random humidity percentages
  windspeed = rnorm(100, mean = 10, sd = 3) # Random wind speeds
)
# Add some outliers
day$temp[c(10, 20, 30)] <- c(40, 45, 50)   # Adding extreme temperatures
day$hum[c(15, 25, 35)] <- c(10, 5, 0)     # Adding extreme humidity values
day$windspeed[c(40, 50, 60)] <- c(25, 30, 35) # Adding extreme wind speeds

# Add some missing values
day$temp[c(5, 15)] <- NA
day$hum[c(10, 20)] <- NA
day$windspeed[c(30, 40)] <- NA

# View the first few rows of the dataset
print(head(day))
print(sum(is.na(day)))
boxplot(day[, c("temp", "hum", "windspeed")], main = "Boxplots of Raw Data")


# Handle outliers in specific columns
for(i in c("hum", "windspeed")) {
  data <- unlist(day[i])
  outliers <- boxplot.stats(data)$out
  data[data %in% outliers] <- NA
  day[i] <- data
}

# Check for missing values after handling outliers
print(sum(is.na(day)))

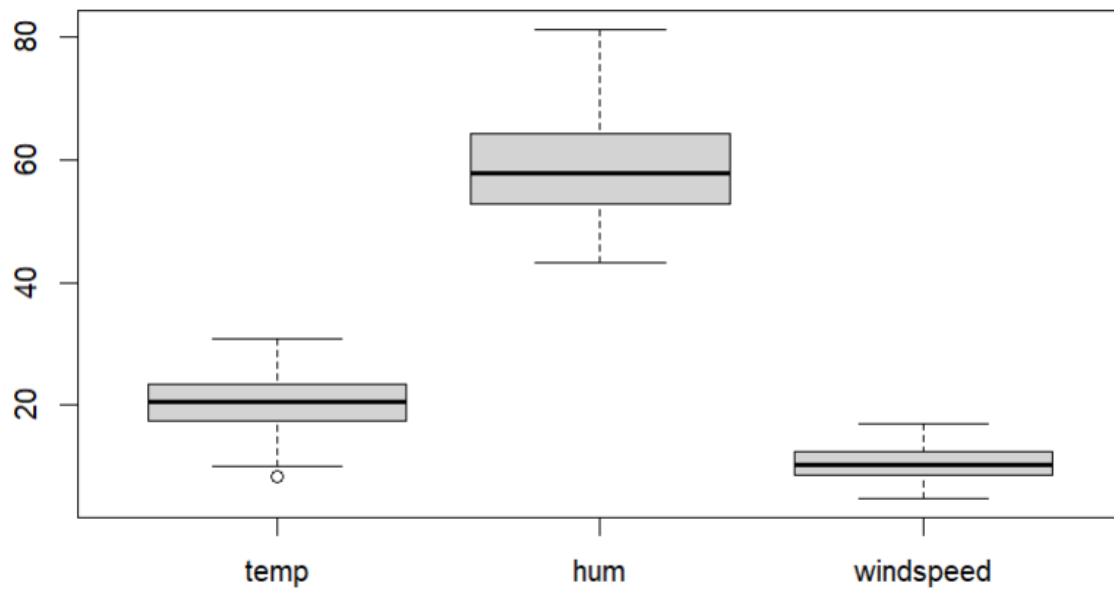
# Drop rows with missing values
day_clean <- drop_na(day)
boxplot(day_clean[, c("temp", "hum", "windspeed")], main = "Boxplots of Cleaned Data")
```

## **Output:**

**Boxplots of Raw Data**



**Boxplots of Cleaned Data**





## 5. Perform analytics on any standard data set

```
# Install and load necessary libraries
if(!require(titanic)) install.packages("titanic")
if(!require(tidyverse)) install.packages("tidyverse")
if(!require(ggcorrplot)) install.packages("ggcorrplot")
library(titanic)
library(tidyverse)
library(ggcorrplot)

# Load Titanic dataset from the titanic package
data("titanic_train")
Titanic_df <- titanic_train

# Convert relevant columns to factors
Titanic_df$Survived <- as.factor(Titanic_df$Survived)
Titanic_df$Pclass <- as.factor(Titanic_df$Pclass)
Titanic_df$Sex <- as.factor(Titanic_df$Sex)

# Remove rows with missing 'Age' values for simplicity
Titanic_df <- Titanic_df %>% drop_na(Age)

# View the first few rows of the dataset
head(Titanic_df)

# Check the structure of the dataset
str(Titanic_df)

# Check for missing values
colSums(is.na(Titanic_df))

# Summary statistics of numerical columns
summary(Titanic_df)

# Summary statistics by survival status
Titanic_df %>% group_by(Survived) %>% summarize(across(where(is.numeric), mean,
```

```
na.rm = TRUE))
```

```
# Visualization
```

```
# 1. Bar plot of survival counts
```

```
ggplot(Titanic_df, aes(x = Survived)) +  
  geom_bar(fill = "skyblue") +  
  labs(title = "Survival Count", x = "Survived", y = "Count")
```

```
# 2. Histogram of Age by Survival
```

```
ggplot(Titanic_df, aes(x = Age, fill = Survived)) +  
  geom_histogram(position = "dodge", bins = 20, alpha = 0.7) +  
  labs(title = "Age Distribution by Survival", x = "Age", y = "Frequency")
```

```
# 3. Boxplot of Fare by Survival
```

```
ggplot(Titanic_df, aes(x = Survived, y = Fare, fill = Survived)) +  
  geom_boxplot() +  
  labs(title = "Fare Distribution by Survival", x = "Survived", y = "Fare")
```

```
# 4. Bar plot of survival rate by class and gender
```

```
ggplot(Titanic_df, aes(x = Pclass, fill = Survived)) +  
  geom_bar(position = "fill") +  
  facet_wrap(~ Sex) +  
  labs(title = "Survival Rate by Class and Gender", x = "Passenger Class", y = "Survival  
Rate")
```

```
# 5. Correlation analysis for numeric columns (Age, Fare, SibSp, Parch)
```

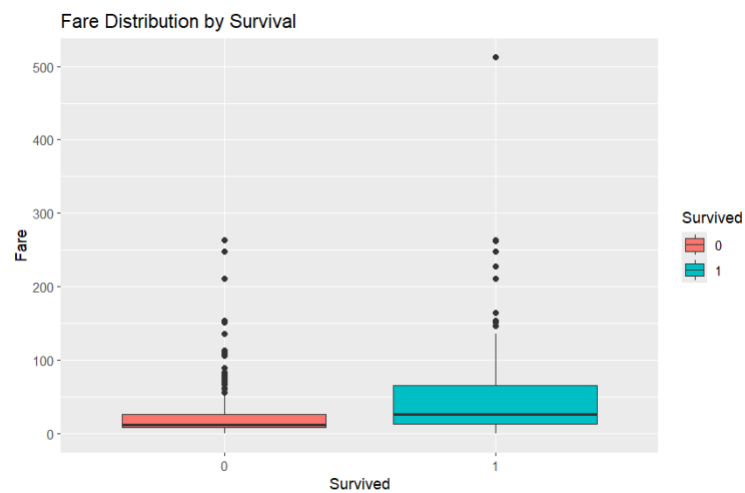
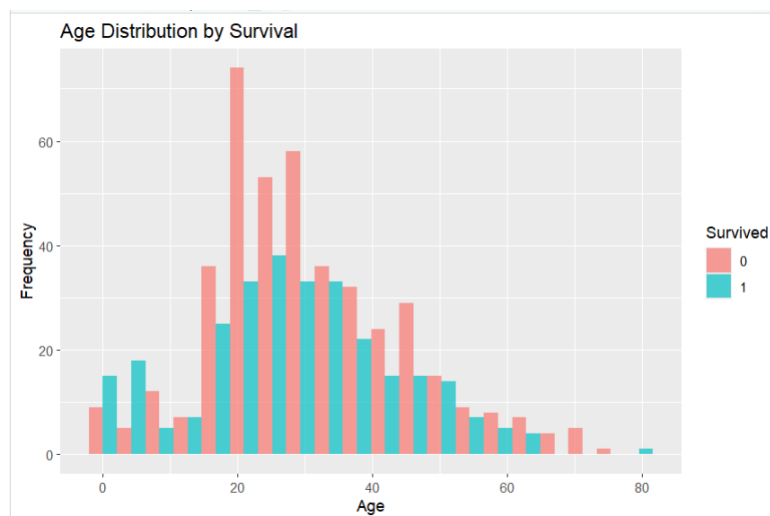
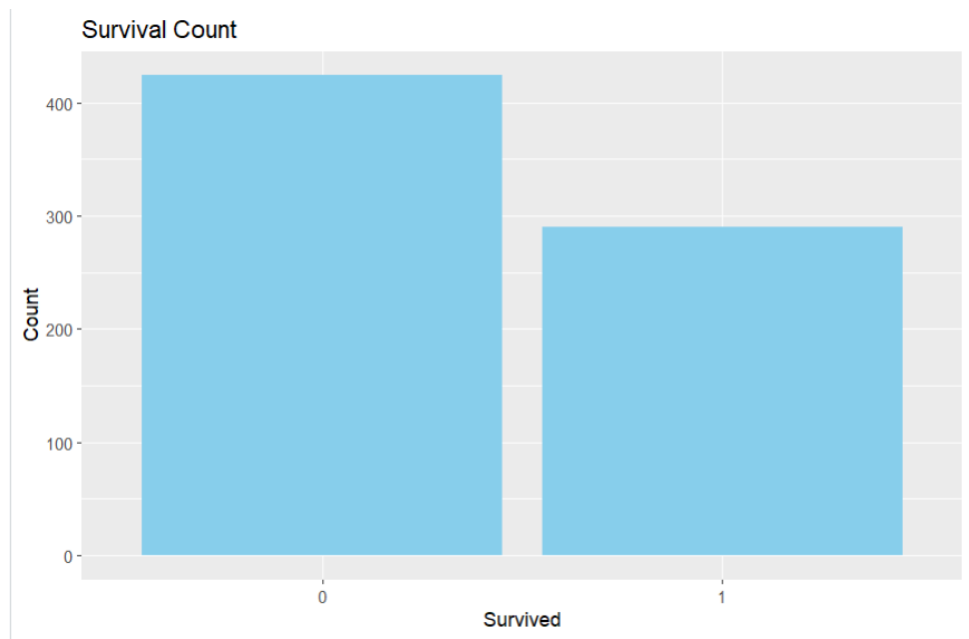
```
numeric_data <- Titanic_df %>% select(Age, Fare, SibSp, Parch)
```



```
cor_matrix <- cor(numeric_data, use = "complete.obs")
```

```
# Plot correlation heatmap
```

```
ggcorrplot(cor_matrix, lab = TRUE)
```

# OUTPUT





## 6. Implement linear regression

```
library(caTools)
library(ggplot2)

# Create the data frame
data <- data.frame(
  Years_Exp = c(1.1, 1.3, 1.5, 2.0, 2.2, 2.9, 3.0, 3.2, 3.2, 3.7),
  Salary = c(39343.00, 46205.00, 37731.00, 43525.00,
            39891.00, 56642.00, 60150.00, 54445.00, 64445.00, 57189.00)
)

# Split the data into training and testing sets
set.seed(123) # Set seed for reproducibility
split = sample.split(data$Salary, SplitRatio = 0.7)
train = subset(data, split == TRUE)
test = subset(data, split == FALSE)

# Fit the linear model
lm.r = lm(formula = Salary ~ Years_Exp, data = train)

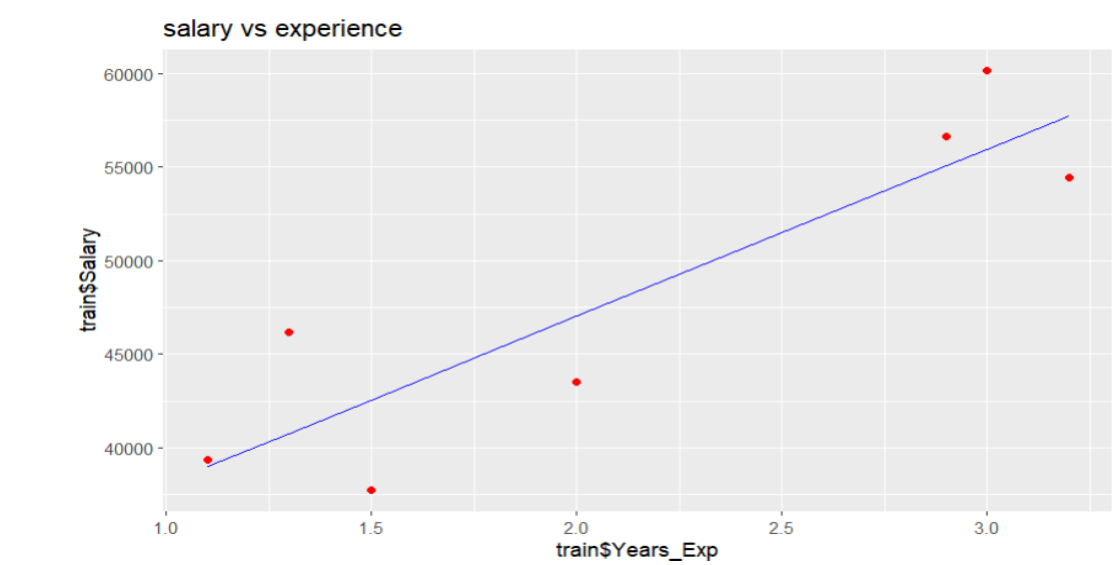
# Print the coefficients
print(coef(lm.r))

# Create the ggplot
ggplot() +
  geom_point(aes(x = train$Years_Exp, y = train$Salary), col = 'red') +
  geom_line(aes(x = train$Years_Exp, y = predict(lm.r, newdata = train)), col = "blue")
+
```

```
ggtitle("Salary vs Experience") +  
xlab("Years of Experience") +  
ylab("Salary")
```

## **Output:**

(Intercept) Years\_Exp  
29172.310 8922.322



## 7. Implement logistic regression

```
library(tidyverse)
library(ROCR)
library(caTools)
library(ggplot2)

# View the mtcars dataset
view(mtcars)

# Split the data into training and testing sets
split <- sample.split(mtcars$vs, SplitRatio = 0.8) # Ensure using 'vs' for splitting
train <- subset(mtcars, split == TRUE) # Use TRUE without quotes
test <- subset(mtcars, split == FALSE) # Use FALSE without quotes

# Build the logistic regression model
logistic_model <- glm(vs ~ wt + disp, data = train, family = binomial)
summary(logistic_model)

# Make predictions
predict_reg <- predict(logistic_model, test, type = "response")

# Convert probabilities to binary outcomes
predict_reg <- ifelse(predict_reg > 0.5, 1, 0)

# Create a confusion matrix
confusion_matrix <- table(test$vs, predict_reg)
print(confusion_matrix)

# Calculate and print classification error and accuracy
missing_classerr <- mean(predict_reg != test$vs)
print(paste("Classification Error = ", missing_classerr))
print(paste("Accuracy = ", (1 - missing_classerr)))

# Plot logistic regression curve
ggplot(train, aes(x = wt + disp, y = vs)) +
  geom_point(alpha = .5) +
  stat_smooth(method = "glm", se = FALSE, method.args = list(family = binomial), col =
"red") +
  labs(title = "Logistic Regression Curve", x = "Weight + Displacement", y = "VS")
```

```

# ROC Curve
ROCPred = prediction(predict_reg, test$vs)
ROCPer = performance(ROCPred, measure = "tpr", x.measure = "fpr")
auc <- performance(ROCPred, measure = "auc")
auc_value <- auc@y.values[[1]]
auc_value <- round(auc_value, 4)

# Plot ROC Curve
plot(ROCPer, colorize = TRUE, print.cutoffs.at = seq(0.1, by = 0.1), main = "ROC
Curve")
abline(a = 0, b = 1)
legend(0.6, 0.4, paste("AUC =", auc_value), title = "AUC", cex = 1)

```

## **Output**

Call:

```
glm(formula = vs ~ wt + disp, family = binomial, data = train)
```

Coefficients:

```

      Estimate Std. Error z value Pr(>|z|)
(Intercept)  2.79114    2.96489   0.941   0.347
wt           0.85989    1.55388   0.553   0.580
disp        -0.02718    0.01456  -1.866   0.062 .

```

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)



Null deviance: 31.841 on 22 degrees of freedom  
Residual deviance: 17.188 on 20 degrees of freedom



AIC: 23.188

Number of Fisher Scoring iterations: 6

```
>
> predict_reg <- predict(logistic_model, test, type = "response")
> predict_reg
      Datsun 710  Hornet Sportabout      Merc 230      Merc 450SLC
      0.864210634    0.017371341    0.841966715    0.189302645
Lincoln Continental  Toyota Corolla  Pontiac Firebird  Porsche 914-2
      0.006385438    0.919574847    0.008373046    0.796023875
      Maserati Bora
      0.089476536
>
> predict_reg <- ifelse(predict_reg > 0.5, 1, 0)
>
> table(test$vs, predict_reg)
predict_reg
0 1
0 5 1
1 0 3
>
> missing_classerr <- mean(predict_reg != test$vs)
> missing_classerr
[1] 0.1111111
> print(paste("accuracy = ", (1 - missing_classerr)))
[1] "accuracy = 0.888888888888889"
>
> library(ggplot2)
>
> #plot logistic regression curve
> ggplot(mtcars, aes(x=wt + disp, y=vs)) +
+   geom_point(alpha=.5) +
+   stat_smooth(method="glm", se=FALSE, method.args = list(family=binomial),
+     col="red")
`geom_smooth()` using formula = 'y ~ x'
> auc
[1] 0.9166667
```



## 8. Construct decision tree for weather data set

```
sample = sample(c(TRUE, FALSE), nrow(weatherdata), replace = TRUE, prob = c(0.8, 0.2))
```

```
train <- weatherdata[sample, ]
```

```
test <- weatherdata[!sample, ]
```

```
library(partykit)
```

```
model <- ctree(RainTomorrow ~ ., train)
```

```
plot(model)
```

```
predict_model <- predict(model, test)
```

```
predict_model
```

```
mat <- table(test$RainTomorrow, predict_model)
```

```
mat
```

```
accuracy <- sum(diag(mat)) / sum(mat)
```

```
accuracy
```

### Output:

```
predict_model
```

```
0.0478723404255319 0.175 0.625
```

```
0      52  10  5
```


```
1       7   4  4
```

```
>
```

```
> accuracy <- sum(diag(mat)) / sum(mat)
```

```
> accuracy
```

```
[1] 0.6829268
```



## 9. Analyse time-series data

```
positiveCases <- c(580, 7813, 28266, 59287, 75700, 87820, 95314, 126214,
218843, 471497, 936851, 1508725, 2072113)
deaths <- c(17, 270, 565, 1261, 2126, 2800,
3285, 4628, 8951, 21283, 47210,
88480, 138475)

library(lubridate)

# output to be created as png file
png(file="multivariateTimeSeries.png")

# creating multivariate time series object
# from date 22 January, 2020
mts <- ts(cbind(positiveCases, deaths),
start = decimal_date(ymd("2020-01-22")),
frequency = 365.25 / 7)

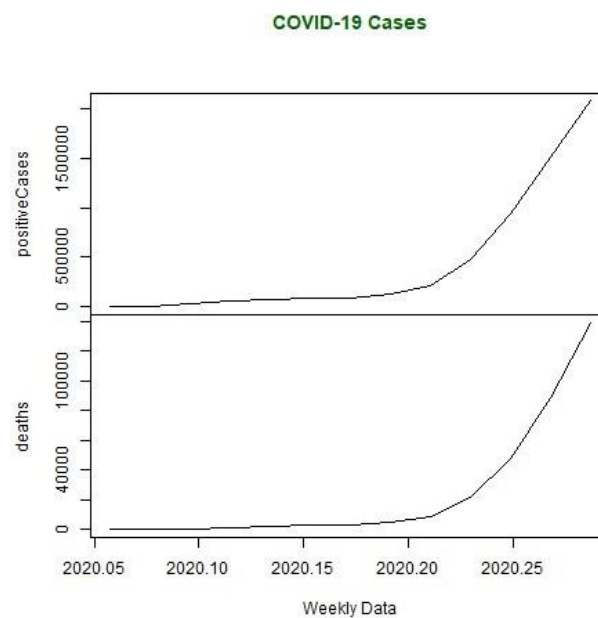
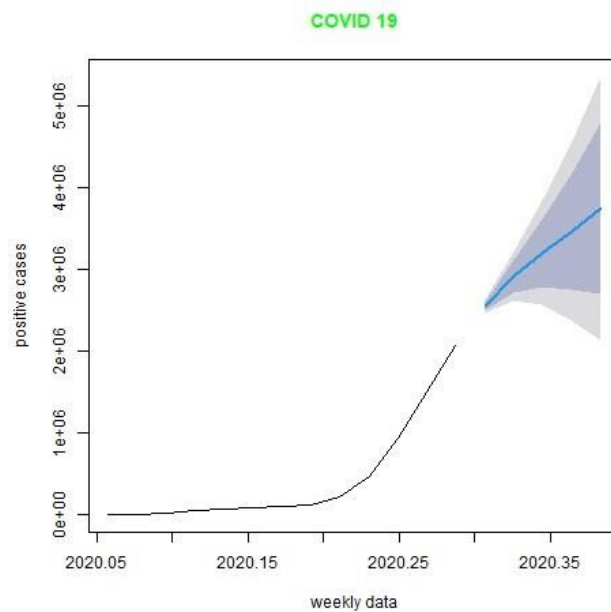
# plotting the graph
plot(mts, xlab = "Weekly Data",
main = "COVID-19 Cases",
col.main = "darkgreen")

library(forecast)
library(lubridate)
png(file = "timeseries.png")
mts1 <- ts(positiveCases, decimal_date(ymd("2020-01-22")), frequency =
365.25/7)
fit <- auto.arima(mts1)
fit <- forecast(fit, 5)
```

```
plot(forecast(fit, 5), xlab="weekly data", ylab = "positive cases", main = "COVID
19", col.main = "green")
```

```
dev.off()
```

## **Output:**



## 10. Work on any data visualization tool

```
view(airquality)

barplot(airquality$Ozone,
        main = 'Ozone Concentration in air',
        xlab = 'ozone levels', horiz = TRUE)

hist(airquality$Temp, main = "La Guardia Airport's\
Maximum Temperature(Daily)",
     xlab = "Temperature(Fahrenheit)",
     xlim = c(50, 125), col = "yellow",
     freq = TRUE)

boxplot(airquality[, 0:4],
        main = 'Box Plots for Air Quality Parameters')

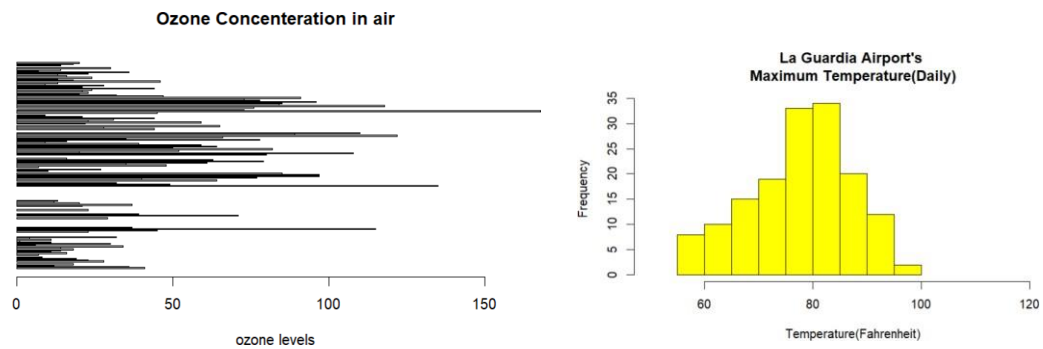
plot(airquality$Ozone, airquality$Month,
     main = "Scatterplot Example",
     xlab = "Ozone Concentration in parts per billion",
     ylab = " Month of observation ", pch = 19)

data <- matrix(rnorm(50, 0, 5), nrow = 5, ncol = 5)


# Column names
colnames(data) <- paste0("col", 1:5)
rownames(data) <- paste0("row", 1:5)

# Draw a heatmap
heatmap(data)
```

# Output



## Box Plots for Air Quality Parameters



## Scatterplot Example

